

# Crash Course in Reinforcement Learning – Cheat Sheet

Basic notation:

- $o \in \mathbb{O}$  – observation (vector of real numbers),
- $s \in \mathbb{S}$  – state,  $o_t = s_t$  for all  $t$ ,
- $a \in \mathbb{A} = \{a^{(1)}, \dots, a^{(A)}\}$  – action,
- $r \in \mathbb{R}$  – reward,
- $\tau = (s_1, a_1, r_1, s_2, \dots, s_T, a_T, r_T, s_{\text{END}})$  – trajectory,
- $R_t = r_t + r_{t+1} + r_{t+2} + \dots + r_T = \sum_{j=t}^T r_j$  – return,
- $R_t^\gamma = \gamma^0 r_t + \gamma^1 r_{t+1} + \gamma^2 r_{t+2} + \dots + \gamma^{T-t} r_T = \sum_{j=t}^T \gamma^{j-t} r_j$  – discounted return,
  - $0 < \gamma < 1$ , usually  $0.9 < \gamma < 1$ ,
  - $\gamma = 1 - \frac{1}{h}$ , where  $h$  is the horizon,
- $R(\tau) = R(s_1, a_1, r_1, s_2, \dots, s_T, a_T, r_T, s_{\text{END}}) = \sum_{j=1}^T r_j$ .

Probabilities:

- $p(s_1)$  – initial state (environment),
- $\pi_\theta(a_t \mid s_t)$  – action (agent),
- $p(r_t, s_{t+1} \mid s_t, a_t)$  – state transition (environment),
- $p(r_t, s_{\text{END}} \mid s_t, a_t)$  – end of episode (environment),
- $\pi_\theta(\tau) = p(s_1) \pi_\theta(a_1 \mid s_1) p(r_1, s_2 \mid s_1, a_1) \dots p(r_T, s_{\text{END}} \mid s_T, a_T)$  – trajectory (agent+environment).

Policy:

$$\begin{aligned}\pi : \mathbb{O} \ni o &\mapsto \left( \pi(a^{(1)} \mid o), \dots, \pi(a^{(A)} \mid o) \right) \in \mathbb{R}^A \\ 0 &\leq \pi(a^{(1)} \mid o), \dots, \pi(a^{(A)} \mid o) \leq 1 \\ \pi(a^{(1)} \mid o) &+ \dots + \pi(a^{(A)} \mid o) = 1\end{aligned}$$

## 1 Policy Gradient

Objective:

$$\begin{aligned}\mathbb{E}_{\tau \sim \pi_\theta} R(\tau) &\approx \frac{1}{N} \sum_{j=1}^N R(\tau_j) \\ \nabla_\theta \mathbb{E}_{\tau \sim \pi_\theta} R(\tau) &= \mathbb{E}_{\tau \sim \pi_\theta} \frac{\nabla_\theta \pi_\theta(\tau)}{\pi_\theta(\tau)} R(\tau) = \mathbb{E}_{\tau \sim \pi_\theta} \sum_{t=1}^T \frac{\nabla_\theta \pi_\theta(a_t \mid s_t)}{\pi_\theta(a_t \mid s_t)} R(\tau) = \mathbb{E}_{\tau \sim \pi_\theta} \sum_{t=1}^T \nabla_\theta \ln \pi_\theta(a_t \mid s_t) \cdot R(\tau) \\ &= \mathbb{E}_{\tau \sim \pi_\theta} \sum_{t=1}^T \nabla_\theta \ln \pi_\theta(a_t \mid s_t) \cdot R_t \approx \mathbb{E}_{\tau \sim \pi_\theta} \sum_{t=1}^T \nabla_\theta \ln \pi_\theta(a_t \mid s_t) \cdot R_t^\gamma \approx \frac{1}{N} \sum_{j=1}^N \sum_{t=1}^T \nabla_\theta \ln \pi_\theta(a_t(\tau_j) \mid s_t(\tau_j)) \cdot R_t^\gamma(\tau_j) \\ &= \nabla_\theta \frac{1}{N} \sum_{j=1}^N \sum_{t=1}^T \ln \pi_\theta(a_t(\tau_j) \mid s_t(\tau_j)) \cdot R_t^\gamma(\tau_j) =: \nabla_\theta J(\theta)\end{aligned}$$

Training loop:

1. Sample  $N$  trajectories  $\tau_1, \dots, \tau_N$ .
2. Calculate the function

$$-J(\theta) = -\frac{1}{N} \sum_{j=1}^N \sum_{t=1}^T \ln \pi_\theta(a_t(\tau_j) \mid s_t(\tau_j)) \cdot R_t^\gamma(\tau_j).$$

3. Ask PyTorch to minimize this function by performing gradient descent step.

## 2 Value Function

Policy gradient with baseline ( $b : \mathbb{S} \rightarrow \mathbb{R}$ ):

$$\mathbb{E}_{\tau \sim \pi_\theta} \sum_{t=1}^T \nabla_\theta \ln \pi_\theta(a_t \mid s_t) \cdot R_t = \mathbb{E}_{\tau \sim \pi_\theta} \sum_{t=1}^T \nabla_\theta \ln \pi_\theta(a_t \mid s_t) \cdot (R_t - b(s_t)).$$

The value function ( $V_{\pi_\theta} : \mathbb{S} \rightarrow \mathbb{R}$ ) approximation with value network ( $V_\psi$ ):

$$V_\psi \approx V_{\pi_\theta}(s) := \mathbb{E}_{\tau \sim \pi_\theta, s_1(\tau)=s} R(\tau) = \mathbb{E}_{\tau \sim \pi_\theta, s_t(\tau)=s} R_t(\tau)$$

Advantage learning (value network as a baseline):

$$\mathbb{E}_{\tau \sim \pi_\theta} \sum_{t=1}^T \nabla_\theta \ln \pi_\theta(a_t \mid s_t) \cdot R_t = \mathbb{E}_{\tau \sim \pi_\theta} \sum_{t=1}^T \nabla_\theta \ln \pi_\theta(a_t \mid s_t) \cdot (R_t - V_\psi(s_t)) = \mathbb{E}_{\tau \sim \pi_\theta} \sum_{t=1}^T \nabla_\theta \ln \pi_\theta(a_t \mid s_t) \cdot A_t$$

Training loop:

1. Sample  $N$  trajectories  $\tau_1, \dots, \tau_N$ .
2. Improve the parameters of our value network by minimizing the loss:

$$\frac{1}{M} \sum_{j=1}^N \sum_{t=1}^T (R_t(\tau_j) - V_\psi(s_t(\tau_j)))^2,$$

where  $M$  is the total number of samples in all trajectories.

3. Calculate the function

$$J(\theta) = \frac{1}{N} \sum_{j=1}^N \sum_{t=1}^T \ln \pi_\theta(a_t(\tau_j) \mid s_t(\tau_j)) \cdot [R_t(\tau_j) - V_\psi(s_t(\tau_j))],$$

4. Ask PyTorch to minimize the function  $-J(\theta)$  by performing gradient descent step.

## 3 PPO

Let's define  $\rho_t$  as:

$$\rho_t := \frac{\pi_\theta(a_t \mid s_t)}{\pi_{\theta\text{OLD}}(a_t \mid s_t)}, \quad \nabla_\theta \rho_t = \frac{\nabla_\theta \pi_\theta(a_t \mid s_t)}{\pi_{\theta\text{OLD}}(a_t \mid s_t)}.$$

As long as for a given  $\epsilon$  (e.g.  $\epsilon = 0.2$ ):

$$1 - \epsilon < \rho_t < 1 + \epsilon,$$

we assume:

$$\begin{aligned} \pi_\theta(s_t) &\approx \pi_{\theta\text{OLD}}(s_t) \\ \mathbb{E}_{\tau \sim \pi_\theta} \frac{\nabla_\theta \pi_\theta(a_t \mid s_t)}{\pi_\theta(a_t \mid s_t)} &\approx \mathbb{E}_{\tau \sim \pi_{\theta\text{OLD}}} \frac{\nabla_\theta \pi_\theta(a_t \mid s_t)}{\pi_{\theta\text{OLD}}(a_t \mid s_t)} \\ \mathbb{E}_{\tau \sim \pi_\theta} \sum_{t=1}^T \nabla_\theta \ln \pi_\theta(a_t \mid s_t) \cdot A_t &\approx \mathbb{E}_{\tau \sim \pi_{\theta\text{OLD}}} \sum_{t=1}^T \nabla_\theta \rho_t \cdot A_t \end{aligned}$$

Surrogate objective:

$$\mathcal{J} := \mathbb{E}_{\tau \sim \pi_{\theta \text{OLD}}} \sum_{t=1}^T \rho_t A_t$$

$$\nabla_{\theta} \mathcal{J} = \nabla_{\theta} \mathbb{E}_{\tau \sim \pi_{\theta \text{OLD}}} \sum_{t=1}^T \rho_t \cdot A_t = \mathbb{E}_{\tau \sim \pi_{\theta \text{OLD}}} \sum_{t=1}^T \nabla_{\theta} \rho_t \cdot A_t$$

Clipped surrogate objective:

$$\tilde{\mathcal{J}} := \mathbb{E}_{\tau \sim \pi_{\theta \text{OLD}}} \sum_{t=1}^T \left[ \min(\rho_t \cdot A_t, \text{clip}(\rho_t, 1 - \epsilon, 1 + \epsilon) \cdot A_t) \right]$$

$$\nabla_{\theta} \tilde{\mathcal{J}} = \mathbb{E}_{\tau \sim \pi_{\theta \text{OLD}}} \sum_{t=1}^T \nabla_{\theta} \left[ \min(\rho_t \cdot A_t, \text{clip}(\rho_t, 1 - \epsilon, 1 + \epsilon) \cdot A_t) \right]$$

Training loop:

1. Sample  $N$  trajectories  $\tau_1, \dots, \tau_N$ .
2. Improve the parameters of our value network by minimizing the loss:

$$\frac{1}{M} \sum_{j=1}^T \sum_{t=1}^N \left[ R_t(\tau_j) - V_{\psi}(s_t(\tau_j)) \right]^2,$$

where  $M$  is the total number of samples in all trajectories.

3. Calculate the function

$$\tilde{J}(\theta) = \frac{1}{M} \sum_{j=1}^N \sum_{t=1}^T \left[ \min(\rho_t \cdot A_t, \text{clip}(\rho_t, 1 - \epsilon, 1 + \epsilon) \cdot A_t) \right],$$

where  $M$  is the batch size.

4. Ask PyTorch to minimize the function  $-\tilde{J}(\theta)$  by performing gradient descent step.